

Turbo Python Performance To Achieve 100x Faster

June 2023

Challenges and Perceptions:

- Python GIL design is slow**
- Python is a slow at runtime**
- Python performance is slower comparing to c++ and java**

Can it be improved?

- Let's find it out**

Techniques used in this performance testing

- Step up and use built-in functions**
- Use vectorization**
- Use math functions**
- Use multi-processing, concurrency**
- Memoization and caching**
- Use different language at server backend**
- Engineering thoughts**

Note:

**Often many things can impact python runtime performance
From hardware, cpu, memory, latency in addition to code
Ideas presented here focus on python source code only**

```
In [1]: 1 # Use built-in functions and libraries, they are tested and optimized
        2 import string
        3 def upper_basic(n):
        4     newList = []
        5     for w in string.ascii_lowercase*n:
        6         newList.append(w.upper())
```

```
In [2]: 1 %timeit upper_basic(1000)
```

4.12 ms \pm 23.1 μ s per loop (mean \pm std. dev. of 7 runs, 100 loops each)

```
In [3]: 1 def upper_o2(n):
        2     newList = map(str.upper, string.ascii_lowercase*n)
```

```
In [72]: 1 %timeit upper_o2(1000)
```

1.22 μ s \pm 18.6 ns per loop (mean \pm std. dev. of 7 runs, 1,000,000 loops each)

60% better when using map() which does elementwise operation

Use vectorization

apply operations to all elements of an array in one go

"for" loop manipulates one row at a time

```
In [20]: 1 def find_sum(n):  
2         total = 0  
3         for i in range(n):  
4             total += i
```

```
In [21]: 1 %timeit find_sum(1_000_000)
```

74.9 ms \pm 1.22 ms per loop (mean \pm std. dev. of 7 runs, 10 loops each)

```
In [22]: 1 import numpy as np  
2         def find_sum_vector(n):  
3             total = 0  
4             total = np.sum(np.arange(n))
```

```
In [23]: 1 %timeit find_sum_vector(1_000_000)
```

2.08 ms \pm 78.9 μ s per loop (mean \pm std. dev. of 7 runs, 100 loops each)

```
In [113]: 1 (74.9-2.08)/74.9
```

```
Out[113]: 0.9722296395193591
```

97% better when using vectorization in numpy

Deep learning multi-linear regression calculations

$$y = m_1x_1 + m_2x_2 + m_3x_3 + m_4x_4 + m_5x_1 + c$$

Use loop for million of rows of calculations is slow

Vectorization is the optimal solution

```
In [38]: 1 # create random data
2 import numpy as np
3 m = np.random.rand(1,5)
4 n = np.random.rand(100000,5)
5 m.shape, n.shape
6
```

```
Out[38]: ((1, 5), (100000, 5))
```

```
In [39]: 1 # use loop for calculations
2 import numpy as np
3
4 def loop_reg_sum(col, row):
5     m = np.random.rand(1,col)
6     n = np.random.rand(row,col)
7     result = []
8     for i in range(row):
9         total = 0
10        for j in range(col):
11            total += n[i][j]*m[0][j]
12        # print(i, total)
13        result.append(total)
14
```

```
In [40]: 1 %timeit loop_reg_sum(5, 100_000)
```

407 ms ± 7.12 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

```
In [41]: 1 # use vectorization
2 def vec_reg_sum(col, row):
3     m = np.random.rand(1,col)
4     n = np.random.rand(row,col)
5     result = np.dot(n, m.T)
6
```

```
In [42]: 1 %timeit vec_reg_sum(5, 100_000)
```

6.45 ms ± 270 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)

407 vs 6.45 Regression using python loop and numpy vectorization

```

In [49]: 1 # use built-in sum
          2 def sum_range(n=1_000_000):
          3     return sum(range(n))

In [50]: 1 # use numpy (implemented in c, faster)
          2 import numpy
          3 def sum_numpy(n=1_000_000):
          4     return numpy.sum(numpy.arange(n))    # this is a one c call, but a whole array is created in memory
          5

In [51]: 1 # use math knowledge
          2 def sum_math(n=1_000_000):
          3     return (n * (n-1)) // 2

In [52]: 1 import timeit
          2
          3 print('while loop\t\t', timeit.timeit(while_loop, number = 1))
          4 print('for loop\t\t', timeit.timeit(for_loop, number = 1))
          5 print('for_loop_with_increment\t\t', timeit.timeit(for_loop_with_increment, number = 1))
          6 print('for_loop_with_test\t\t', timeit.timeit(for_loop_with_test, number = 1))
          7 print('for_loop_with_increment_and_test\t\t', timeit.timeit(for_loop_with_increment_and_test, number = 1))
          8 print('sum_range\t\t', timeit.timeit(sum_range, number = 1))
          9 print('sum_numpy\t\t', timeit.timeit(sum_numpy, number = 1))
         10 print('sum_math\t\t', timeit.timeit(sum_math, number = 1))
         11
         12 # python programming consideration
         13 # use math formula
         14 # use c implementation
         15 # use built-in function, sum, map ... which loops for you
         16 # for or while loop
         17
         18

while loop          0.11545043099999993
for loop            3.713000069183181e-06
for_loop_with_increment    0.11514549399998941
for_loop_with_test        0.0975002409999206
for_loop_with_increment_and_test    0.1432448260000001
sum_range           0.022303372999999738
sum_numpy           0.005475752999927863
sum_math            2.3169999394667684e-06

```

Different loops and their performances

```
In [53]: 1 ## memoization or cache to optimize
2 # useful for recursive functions, or operations used over and over again
3 # you don't want to repeat to calculate values again
```

```
In [115]: 1 # use cache dict
2 from time import perf_counter
3 from functools import wraps
4
5 def memoize(func):
6     cache = {}
7
8     @wraps(func)
9     def wrapper(*args, **kwargs):
10         key = str(args) + str(kwargs)
11         if key not in cache:
12             cache[key] = func(*args, **kwargs)
13         return cache[key]
14     return wrapper
15
```

```
In [116]: 1 # fibonacci using memoize
2 def fibonacci_plain(n=100) -> int:
3     if n < 2:
4         return n
5     return fibonacci_plain(n-1) + fibonacci_plain(n-2)
```

```
In [*]: 1 # no memoization call, very slow, cpu humming, 20 mins still running, killed this cell
2 start = perf_counter()
3 fibonacci_plain()
4 end = perf_counter()
5 print(end-start)
```

```
In [57]: 1 # fibonacci using memoize
2 @memoize
3 def fibonacci(n=1000) -> int:
4     if n < 2:
5         return n
6     return fibonacci(n-1) + fibonacci(n-2)
```

```
In [58]: 1 # get result instantly
2 print('fibonacci with memorize\t\t', timeit.timeit(fibonacci, number = 1))
```

fibonacci with memorize 0.0036856550000265997

Memoization and caching reducing intermediate operations

```

27 outputs = []
28 for url in urls.values():
29     print(url)
30     outputs = outputs + [requests.get(url).text]
31     #print(outputs)
32
33 count_https = []
34 count_http = []
35 for output in outputs:
36     count_https += re.findall("https://", output)
37     count_http += re.findall("http://", output)
38
39 print(len(count_https), len(count_http))
40
41 # index = 0
42 # while count_https[index]:
43 #     if index >= len(count_https):
44 #         break
45 #     index += 1
46
47 start = time.perf_counter()
48 count_words_in_web_page()
49 elapsed = time.perf_counter() - start
50 print(f'{elapsed:.2f} seconds')
51

```

```

https://google.com
https://yahoo.com
https://microsoft.com
https://google.com
https://apple.com
https://ibm.com
https://amazon.com
https://twitter.com
https://tiktok.com
https://oracle.com
https://intel.com
https://tesla.com
https://nasa.com
https://ebay.com
https://wikipedia.com
3071 732
10.50 seconds

```

```

14 "2": "https://yahoo.com",
15 "3": "https://microsoft.com",
16 "4": "https://google.com",
17 "5": "https://apple.com",
18 "6": "https://ibm.com",
19 "7": "https://amazon.com",
20 "8": "https://twitter.com",
21 "9": "https://tiktok.com",
22 "10": "https://oracle.com",
23 "11": "https://intel.com",
24 "12": "https://tesla.com",
25 "13": "https://nasa.com",
26 "14": "https://ebay.com",
27 "15": "https://wikipedia.com"
28 }
29
30 # mark as async
31 async def count_words_in_web_page_async():
32     outputs = []
33
34     async with httpx.AsyncClient() as client:
35         tasks = (client.get(url) for url in urls.values())
36         reqs = await asyncio.gather(*tasks) # waits for task, but await till all donee
37
38         outputs = [req.text for req in reqs]
39         #print(outputs)
40
41     count_https, count_http = [], []
42     for output in outputs:
43         count_https += re.findall("https://", output) # text processing, not use pre-compiled re
44         count_http += re.findall("http://", output)
45     # print(count_https)
46     # print(count_http)
47
48 start = time.perf_counter()
49 await count_words_in_web_page_async() # schedule func to run
50 # asyncio.run(count_words_in_web_page_async()) # for python>3.7 and ipython < 7.0
51 elapsed = time.perf_counter() - start
52 print(f'{elapsed:.2f} seconds')
53
54
55

```

1.03 seconds

10.50 vs 1.03 - Use async for web text scraping


```

(base) user-2:bin user$ cat main.rs
extern crate webserver;

use webserver::ThreadPool;
use std::net::TcpListener;
use std::io::prelude::*;
use std::net::TcpStream;
use std::fs::File;
use std::thread;
use std::time::Duration;

fn main() {

    let listener = TcpListener::bind("127.0.0.1:7878").unwrap();
    let pool = ThreadPool::new(8);

    for stream in listener.incoming() {

        let stream = stream.unwrap();

        pool.execute(|| {
            handle_connection(stream);
        });

    }

}

fn handle_connection(mut stream: TcpStream) {

    let mut buffer = [0; 512];

    stream.read(&mut buffer).unwrap();

    let get = b"GET / HTTP/1.1\r\n";
    let sleep = b"GET /sleep HTTP/1.1\r\n";

    let (status_line, filename) = if buffer.starts_with(get) {
        ("HTTP/1.1 200 OK\r\n\r\n", "hello.html")
    } else if buffer.starts_with(sleep) {
        thread::sleep(Duration::from_secs(5));
        ("HTTP/1.1 200 OK\r\n\r\n", "hello.html")
    } else {
        ("HTTP/1.1 404 NOT FOUND\r\n\r\n", "404.html")
    };

    let mut file = File::open(filename).unwrap();
    let mut contents = String::new();

    file.read_to_string(&mut contents).unwrap();

    let response = format!("{}", status_line, contents);

    stream.write(response.as_bytes()).unwrap();
    stream.flush().unwrap();

}

```

Rust multi thread server

Python web api call

```
url = 'http://localhost:7878/'

def fetch(session, url):
    with session.get(url) as response:
        #print(response)
        pass

@timer(1, 1)
def main():
    with requests.Session() as session:
        for _ in range(5000):
            fetch(session, url)

import requests
import timeit
from multiprocessing.pool import Pool

url = 'http://localhost:7878/'

def fetch(session, url):
    with session.get(url) as response:
        #print(response)
        pass

def timer(number, repeat):
    def wrapper(func):
        runs = timeit.repeat(func, number=number, repeat=repeat)
        print(sum(runs) / len(runs))
    return wrapper

if __name__ == "__main__":

    @timer(1, 1)
    def task():
        with Pool() as pool:
            with requests.Session() as session:
                pool.starmap(fetch, [(session, url) for _ in range(5000)])

(base) user-2:client user$
(base) user-2:client user$ cat ../readme.txt
# start server
(base) user-2:rust_web_server_concurrent user$ cargo run
    Finished dev [unoptimized + debuginfo] target(s) in 0.00s
    Running `target/debug/main`

# start 01_simple-http.py
# synchronous calls
(base) user-2:client user$ python 01_simple_http_sync.py
11.093317224

# multiprocessing using multi-cores to run multi processes
(base) user-2:client user$ python 02_multi_processing_http.py
2.7904895300000003
```

Python web api call

```
url = 'http://localhost:7878/'

def fetch(session, url):
    with session.get(url) as response:
        #print(response)
        pass

@timer(1, 1)
def main():
    with requests.Session() as session:
        for _ in range(5000):
            fetch(session, url)

import requests
import timeit
from multiprocessing.pool import Pool

url = 'http://localhost:7878/'

def fetch(session, url):
    with session.get(url) as response:
        #print(response)
        pass

def timer(number, repeat):
    def wrapper(func):
        runs = timeit.repeat(func, number=number, repeat=repeat)
        print(sum(runs) / len(runs))
    return wrapper

if __name__ == "__main__":

    @timer(1, 1)
    def task():
        with Pool() as pool:
            with requests.Session() as session:
                pool.starmap(fetch, [(session, url) for _ in range(5000)])

(base) user-2:client user$
(base) user-2:client user$ cat ../readme.txt
# start server
(base) user-2:rust_web_server_concurrent user$ cargo run
    Finished dev [unoptimized + debuginfo] target(s) in 0.00s
    Running `target/debug/main`

# start 01_simple-http.py
# synchronous calls
(base) user-2:client user$ python 01_simple_http_sync.py
11.093317224

# multiprocessing using multi-cores to run multi processes
(base) user-2:client user$ python 02_multi_processing_http.py
2.7904895300000003
```

```
In [2]: 1 !python -V
```

Python 3.10.11

```
In [1]: 1 import math
        2 import numpy as np
```

Factorial goes faster

```
In [39]: 1 # basic factorial
        2 %timeit math.prod(range(1, 150))
```

10.6 μ s \pm 32.3 ns per loop (mean \pm std. dev. of 7 runs, 100,000 loops each)

```
In [33]: 1 math.prod(range(1,6))
```

Out[33]: 120

```
In [34]: 1 def f(x):
        2     return x * f(x-1) if x > 1 else 1
```

```
In [28]: 1 f(5)
```

Out[28]: 120

```
In [40]: 1 %timeit math.factorial(150)
```

1.8 μ s \pm 3.93 ns per loop (mean \pm std. dev. of 7 runs, 1,000,000 loops each)

```
In [44]: 1 (1.8/10.6)
```

Out[44]: 0.169811320754717

Use new versions

compute factorials like a boss

Winning ideas:

Left shift is cheaper than multiplying by two

Pulling out events leaves recurring odd factories

Dynamic programming reuses previously computed odd factorials

```
In [75]: 1 s = ""
2 so = ""
3 se = ""
4 for i in range(1,20):
5     s += str(i) + '*'
6     if i%2 == 0:
7         se += str(i) + '*'
8     else:
9         so += str(i) + '*'
10    print(s, so, se)

1*2*3*4*5*6*7*8*9*10*11*12*13*14*15*16*17*18*19* 1*3*5*7*9*11*13*15*17*19* 2*4*6*8*10*12*14*16*18*
```

```
In [71]: 1 1*2*3*4*5*6*7*8*9*10*11*12*13*14*15*16*17*18*19
```

Out[71]: 121645100408832000

```
In [79]: 1 # collect even terms, next divide each even by 2
2 1*3*5*7*9*11*13*15*17*19 *2*4*6*8*10*12*14*16*18
```

Out[79]: 121645100408832000

```
In [81]: 1 # divide terms by two and replace with left shift
2 1*3*5*7*9*11*13*15*17*19 *1*2*3*4*5*6*7*8*9 << 9
```

Out[81]: 121645100408832000

```
In [84]: 1 1*3*5*7*9*11*13*15*17*19 *1*3*5*7*9 * 1*2*3*4 <<13
```

Out[84]: 121645100408832000

```
In [86]: 1 1*3*5*7*9*11*13*15*17*19 *1*3*5*7*9 * 1*3 *1*2 <<15
```

Out[86]: 121645100408832000

```
In [88]: 1 # replace even term with left shift
2 1*3*5*7*9*11*13*15*17*19 *1*3*5*7*9 * 1*3 <<16
```

Out[88]: 121645100408832000

```
In [93]: 1 # factor-out common subsequences and replace with powers
2 (1*3)**3 *(5*7*9)**2 *(11*13*15*17*19)**1 <<16
```

Out[93]: 121645100408832000

How python 3 implement factorial() performance

Conclusions

- Python is the programming language to solve all problems**
- Python is a popular language for many use cases**
- Performance can be improved**
- Python is improving, by itself and community**

- In the era of ML/AI, no doubt python will become popular**
- To learn, experiment, build quickly MVP in less time**

Q & A